

SCIKIT-LEARN : DATA COMPRESSION VIA DIMENSIONALITY REDUCTION I - PRINCIPAL COMPONENT ANALYSIS (PCA)



(<http://www.addthis.com/bookmark.php?v=250&username=khhong7>)

bogotobogo.com site search:

Principal component analysis (PCA)


Principal component analysis (PCA) is a statistical procedure that uses an orthogonal transformation to convert a set of observations of possibly correlated variables into a set of values of linearly uncorrelated variables called principal components. - wiki (https://en.wikipedia.org/wiki/Principal_component_analysis)

PCA tries to find the directions of maximum variance (direction of orthogonal axes / principal components) in data and projects it onto a new subspace with lower dimension than the original one.



K Hong

google.com/+KHongSanFrancisco



4,892 followers


Ph.D. / Golden Gate Ave, San Francisco / Seoul National Univ / Carnegie Mellon / UC Berkeley / DevOps / Deep Learning / Visualization

Sponsor Open Source development activities and free contents for everyone.



Thank you.

- K Hong (http://bogotobogo.com/about_us.php)



Principal Components Analysis (PCA) (<https://onlinecourses.science.psu.edu/stat857/node/35>)

Here, x_1x_1 and x_2x_2 are the original feature axes, and z_1z_1 and z_2z_2 are the principal components.

Dimensionality reduction via principal component analysis

In order to reduce dimensionality using PCA, we construct a transformation matrix WW which has $d \times kd \times k$ -dimension.

With the WW matrix we can map a sample vector xx onto a new kk -dimensional feature subspace that has fewer dimensions than the original dd -dimensional feature space:

$$\begin{aligned} \mathbf{x} &= [x_1, x_2, \dots, x_d], & \mathbf{x} \in \mathbb{R}^d \\ \mathbf{x} &= [x_1, x_2, \dots, x_d], & \mathbf{x} \in \mathbb{R}^d \\ \downarrow \mathbf{xW}, & & W \in \mathbb{R}^{d \times k} \\ \downarrow \mathbf{xW}, & & W \in \mathbb{R}^{d \times k} \\ \mathbf{z} &= [z_1, z_2, \dots, z_k], & \mathbf{z} \in \mathbb{R}^k \\ \mathbf{z} &= [z_1, z_2, \dots, z_k], & \mathbf{z} \in \mathbb{R}^k \end{aligned}$$

After the transformation from the original dd -dimensional data onto this new kk -dimensional subspace ($k \leq dk \leq d$), the first principal component will have the largest possible variance, and all consequent principal components will have the largest possible variance given that they are uncorrelated (orthogonal) to the other principal components.

Machine Learning with scikit-learn

scikit-learn installation
(/python/scikit-learn/scikit-learn_install.php)

scikit-learn : Features and feature extraction - iris dataset
(/python/scikit-learn/scikit_machine_learning_fe)

scikit-learn : Machine Learning Quick Preview (/python/scikit-learn/scikit_machine_learning_q)

scikit-learn : Data Preprocessing I - Missing / Categorical data)
(/python/scikit-learn/scikit_machine_learning_D_Missing-Data-Categorical-Data.php)

scikit-learn : Data Preprocessing II - Partitioning a dataset / Feature scaling / Feature Selection / Regularization (/python/scikit-learn/scikit_machine_learning_D_II-Datasets-Partitioning-Feature-scaling-Feature-Selection-Regularization.php)

scikit-learn : Data Preprocessing III - Dimensionality reduction vis Sequential feature selection / Assessing feature importance via random forests
(/python/scikit-learn/scikit_machine_learning_D_III-Dimensionality-reduction-via-Sequential-feature-selection-Assessing-feature-importance-via-random-forests.php)

Note that the PCA directions are highly sensitive to data scaling, and most likely we need to standardize the features prior to PCA if the features were measured on different scales and we want to assign equal importance to all features.

Here are the steps of PCA algorithm for dimensionality reduction:

1. Standardize the d -dimensional dataset.
2. Construct the covariance matrix.
3. Decompose the covariance matrix into its eigenvectors and eigenvalues.
4. Select k eigenvectors that correspond to the k largest eigenvalues, where k is the dimensionality of the new feature subspace ($k \leq d$).
5. Construct a projection matrix W from the "top" k eigenvectors.
6. Transform the d -dimensional input dataset \mathbf{x} using the projection matrix W to obtain the new k -dimensional feature subspace.

Eigenvalues & eigenvectors

Continued from the previous section for principal component analysis, in this section we'll standardize the data, construct the covariance matrix, obtain the eigenvalues and eigenvectors of the covariance matrix, and sort the eigenvalues by decreasing order to rank the eigenvectors.

Let's start by loading the Wine dataset from "https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wine.data":

```
%pylab inline
Populating the interactive namespace from numpy and matplotlib

import pandas as pd

df_wine = pd.read_csv(
    'https://archive.ics.uci.edu/ml/machine-learning-databases/wine/wi
    header=None)

df_wine.head()
```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	2.29	5.64	1.04	3.92	1065
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	4.38	1.05	3.40	1050
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	2.81	5.68	1.03	3.17	1185
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	2.18	7.80	0.86	3.45	1480
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	1.82	4.32	1.04	2.93	735

Note that PCA is an **unsupervised** method, which means that information about the class labels is ignored. It shows clear contrast compared with a random forest which uses the class membership information to compute the node impurities, variance measures the spread of values along a feature axis. Recall we did the following for random forest:

```
df_wine.columns = ['Class label', 'Alcohol', 'Malic acid', 'Ash',
    'Alcalinity of ash', 'Magnesium', 'Total phenols', 'Flavanoids',
    'Nonflavanoid phenols', 'Proanthocyanins', 'Color intensity', 'Hue',
    'OD280/OD315 of diluted wines', 'Proline']
```

After the read-in, we process the Wine data into separate training (70%) and test (30%) sets and then standardize it to unit variance:

scikit-learn : Data Compression via Dimensionality Reduction I - Principal component analysis (PCA) (/python/scikit-learn/scikit_machine_learning_D_PCA.php)

scikit-learn : Data Compression via Dimensionality Reduction II - Linear Discriminant Analysis (LDA) (/python/scikit-learn/scikit_machine_learning_D

scikit-learn : Data Compression via Dimensionality Reduction III - Nonlinear mappings via kernel principal component (KPCA) analysis (/python/scikit-learn/scikit_machine_learning_D nonlinear-mappings-via-kernel-principal-component-analysis.php)

scikit-learn : Logistic Regression, Overfitting & regularization (/python/scikit-learn/scikit-learn_logistic_regression.php)

scikit-learn : Supervised Learning & Unsupervised Learning - e.g. Unsupervised PCA dimensionality reduction with iris dataset (/python/scikit-learn/scikit_machine_learning_S

scikit-learn : Unsupervised_Learning - KMeans clustering with iris dataset (/python/scikit-learn/scikit_machine_learning_U

scikit-learn : Linearly Separable Data - Linear Model & (Gaussian) radial basis function kernel (RBF kernel) (/python/scikit-learn/scikit_machine_learning_Li

scikit-learn : Decision Tree Learning I - Entropy, Gini, and Information Gain (/python/scikit-learn/scikt_machine_learning_De

scikit-learn : Decision Tree


```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
X, y = df_wine.iloc[:, 1:].values, df_wine.iloc[:, 0].values
X_train, X_test, y_train, y_test = \
    train_test_split(X, y, test_size=0.3, random_state=0)
sc = StandardScaler()
X_train_std = sc.fit_transform(X_train)
X_test_std = sc.fit_transform(X_test)
```

Now we want to construct the covariance matrix which is symmetric with $d \times d$ -dimension, where d is the dataset dimension. The covariance matrix stores the pairwise covariances between the different features.

The covariance between two features x_j and x_k on the population level can be calculated via the equation below:

$$\sigma_{jk} = \frac{1}{N} \sum_{i=1}^N (x_j^i - \mu_j)(x_k^i - \mu_k)$$
$$\sigma_{jk} = \frac{1}{N} \sum_{i=1}^N (x_j^i - \mu_j)(x_k^i - \mu_k)$$

where μ_j and μ_k are the sample means of feature j and k , respectively.

Note that the sample means are zero if we standardize the dataset.

A positive covariance between two features indicates that the features increase or decrease together, while a negative covariance means that the features vary in opposite directions.

A covariance matrix of three features can then be written as A :

$$A = \begin{bmatrix} \sigma_{11}^2 & \sigma_{12} & \sigma_{13} \\ \sigma_{21} & \sigma_{22}^2 & \sigma_{23} \\ \sigma_{32} & \sigma_{32} & \sigma_{33}^2 \end{bmatrix}$$
$$A = \begin{bmatrix} \sigma_{11}^2 & \sigma_{12} & \sigma_{13} \\ \sigma_{21} & \sigma_{22}^2 & \sigma_{23} \\ \sigma_{32} & \sigma_{32} & \sigma_{33}^2 \end{bmatrix}$$

The eigenvectors of the covariance matrix represent the principal components, while the corresponding eigenvalues will define their magnitude.

In the case of the Wine dataset, we can obtain 13 eigenvectors and eigenvalues from the 13×13 covariance matrix.

An eigenvector ν satisfies the following condition where λ is the eigenvalue:

$$A\nu = \lambda\nu$$
$$A\nu = \lambda\nu$$

Learning II - Constructing the Decision Tree (/python/scikit-learn/scikit_machine_learning_C

scikit-learn : Random Decision Forests Classification (/python/scikit-learn/scikit_machine_learning_R

scikit-learn : k-Nearest Neighbors (k-NN) Algorithm (/python/scikit-learn/scikit_machine_learning_k-NN_k-nearest-neighbors-algorithm.php)

scikit-learn : Support Vector Machines (SVM) (/python/scikit-learn/scikit_machine_learning_S

scikit-learn : Support Vector Machines (SVM) II (/python/scikit-learn/scikit_machine_learning_S

Flask with Embedded Machine Learning I : Serializing with pickle and DB setup (/python/Flask/Python_Flask_Em

Flask with Embedded Machine Learning II : Basic Flask App (/python/Flask/Python_Flask_Em

Flask with Embedded Machine Learning III : Embedding Classifier (/python/Flask/Python_Flask_Em

Flask with Embedded Machine Learning IV : Deploy (/python/Flask/Python_Flask_Em

Flask with Embedded Machine Learning V : Updating the classifier (/python/Flask/Python_Flask_Em

scikit-learn : Sample of a spam comment filter using SVM - classifying a good one or a bad one (/python/scikit-learn/scikit_learn_Support_Vecto

We're going to use the **linalg.eig** function from NumPy to obtain the eigenpairs of the Wine covariance matrix:

```
covariant_matrix = np.cov(X_train_std.T)

covariant_matrix[0::5]

array([[ 1.00813008,   0.08797701,   0.23066952,  -0.32868099,   0.2141631 ,
         0.35576761,   0.2991246 ,  -0.16913744,   0.09649074,   0.56962271,
        -0.04781543,   0.07403492,   0.63277882],
       [ 0.35576761,  -0.30124242,   0.12235533,  -0.37018442,   0.16513295,
         1.00813008,   0.88119961,  -0.45396901,   0.6196806 ,  -0.06935051,
         0.45718802,   0.72214462,   0.56326772],
       [-0.04781543,  -0.54992807,  -0.10928021,  -0.25313262,   0.05792599,
         0.45718802,   0.58331869,  -0.3178224 ,   0.32282167,  -0.52395358,
         1.00813008,   0.60022569,   0.2452794 ]])

eigen_values, eigen_vectors = np.linalg.eig(covariant_matrix)

eigen_values, eigen_vectors[:, :5]

(array([ 4.8923083 ,   2.46635032,   1.42809973,   1.01233462,   0.84906459,
         0.60181514,   0.52251546,   0.08414846,   0.33051429,   0.29595018,
         0.16831254,   0.21432212,   0.2399553 ]),
 array([[ 1.46698114e-01,   5.04170789e-01,  -1.17235150e-01,
         2.06254611e-01,  -1.87815947e-01,  -1.48851318e-01,
        -1.79263662e-01,  -5.54687162e-02,  -4.03054922e-01,
        -4.17197583e-01,   2.75660860e-01,   4.03567189e-01,
         4.13320786e-04],
       [ 3.89344551e-01,   9.36399132e-02,   1.80804417e-01,
         1.93179478e-01,   1.40645426e-01,   1.22248798e-02,
         5.31455344e-02,  -4.21265116e-01,   1.35111456e-01,
        -2.80985650e-01,   2.83897644e-01,  -6.18600153e-01,
         9.45645138e-02],
       [ 3.00325353e-01,  -2.79243218e-01,   9.32387182e-02,
         2.41740256e-02,  -3.72610811e-01,   2.16515349e-01,
        -3.84654748e-01,  -1.05383688e-01,  -5.17259438e-01,
        1.97814118e-01,  -1.98844532e-01,  -2.00456386e-01,
        -3.02254353e-01]]))
```

We computed the covariance matrix of the standardized training dataset using the **numpy.cov()** function.

Using the `linalg.eig` function, we performed the eigendecomposition that yielded 13 **eigenvalues** and the corresponding **eigenvectors** stored as columns in a 13×13 matrix.

Since we want to reduce the dimensionality of our dataset by compressing it onto a new feature subspace, we only select the subset of the eigenvectors (principal components) that contains most of the information (variance).

Since the eigenvalues define the magnitude of the eigenvectors, we have to sort the eigenvalues by decreasing magnitude, and we are interested in the top k eigenvectors based on the values of their corresponding eigenvalues.

But before we collect those k most informative eigenvectors, let's plot the variance explained ratios of the eigenvalues.

The variance explained ratio of an eigenvalue λ_j is simply the fraction of an eigenvalue λ_j and the total sum of the eigenvalues:

$$\frac{\lambda_j}{\sum_{j=1}^d \lambda_j}$$
$$\frac{\lambda_j}{\sum_{j=1}^d \lambda_j}$$

Using the NumPy **cumsum()** function, we can then calculate the **cumulative sum** of explained variances, which we will plot via matplotlib's **step()** function:

MACHINE LEARNING ALGORITHMS

Batch gradient descent algorithm (/python/python_numpy_batch_

Single Layer Neural Network - Perceptron model on the Iris dataset using Heaviside step activation function (/python/scikit-learn/Perceptron_Model_with_Ir

Batch gradient descent versus stochastic gradient descent (SGD) (/python/scikit-learn/scikit-learn_batch-gradient-descent-versus-stochastic-gradient-descent.php)

Single Layer Neural Network - Adaptive Linear Neuron using linear (identity) activation function with batch gradient descent method (/python/scikit-learn/Single-Layer-Neural-Network-Adaptive-Linear-Neuron.php)

Single Layer Neural Network : Adaptive Linear Neuron using linear (identity) activation function with stochastic gradient descent (SGD) (/python/scikit-learn/Single-Layer-Neural-Network-Adaptive-Linear-Neuron-with-Stochastic-Gradient-Descent.php)

VC (Vapnik-Chervonenkis) Dimension and Shatter (/python/scikit-learn/scikit_machine_learning_V


Bias-variance tradeoff (/python/scikit-learn/scikit_machine_learning_B variance-Tradeoff.php)

Logistic Regression (/python/scikit-learn/logistic_regression.php)

Maximum Likelihood Estimation (MLE)


```
tot = sum(eigen_values)
var_exp = [(i / tot) for i in sorted(eigen_values, reverse=True)]
cum_var_exp = np.cumsum(var_exp)

bar(range(1,14), var_exp, alpha=0.5, align='center',
      label='individual explained variance')
step(range(1,14), cum_var_exp, where='mid',
      label='cumulative explained variance')
ylabel('Explained variance ratio')
xlabel('Principal components')
legend(loc='best')
show()
```



The plot shows that the first principal component alone accounts for 40 percent of the variance. Also, we can see that the first two principal components combined explain almost 60 percent of the variance in the data.

Transform dataset onto a new principal axes

Now that we have decomposed the covariance matrix into eigen-pairs, we want to transform the Wine dataset onto the new principal component axes.

We're going to sort the eigen-pairs by descending order of the eigenvalues, and construct a projection matrix from the selected eigenvectors. Then, using the projection matrix we will transform the data onto the lower-dimensional subspace.

Let's start by sorting the eigen-pairs by decreasing order of the eigenvalues:

```
eigen_pairs = \
[(np.abs(eigen_values[i]),eigen_vectors[:,i]) for i in range(len(eigen_values))]
eigen_pairs.sort(reverse=True)

eigen_pairs[:5]
```

```
[(4.8923083032737509,
 array([ 0.14669811, -0.24224554, -0.02993442, -0.25519002,  0.12079772,
        0.38934455,  0.42326486, -0.30634956,  0.30572219, -0.09869191,
        0.30032535,  0.36821154,  0.29259713])),
 (2.4663503157592306,
 array([ 0.50417079,  0.24216889,  0.28698484, -0.06468718,  0.22995385,
        0.09363991,  0.01088622,  0.01870216,  0.03040352,  0.54527081,
        -0.27924322, -0.174365  ,  0.36315461])),
 (1.4280997275048455,
 array([-0.11723515,  0.14994658,  0.65639439,  0.58428234,  0.08226275,
        0.18080442,  0.14295933,  0.17223475,  0.1583621 , -0.14242171,
        0.09323872,  0.19607741, -0.09731711])),
 (1.0123346209044966,
 array([ 0.20625461,  0.1304893 ,  0.01515363, -0.09042209, -0.83912835,
        0.19317948,  0.14045955,  0.33733262, -0.1147529 ,  0.07878571,
        0.02417403,  0.18402864,  0.05676778])),
 (0.8490645933450266,
 array([-0.18781595,  0.56863978, -0.29920943, -0.04124995, -0.02719713,
        0.14064543,  0.09268665, -0.08584168,  0.56510524,  0.01323461,
        -0.37261081,  0.08937967, -0.21752948]))]
```

Next, we collect the two eigenvectors that correspond to the two largest values to capture about 60 percent of the variance in this dataset.

(/python/scikit-learn/Maximum-Likelyhood-Estimation-MLE.php)

Neural Networks with backpropagation for XOR using one hidden layer (/python/python_Neural_Networks/Neural_Networks-XOR-Backpropagation.php)

minHash (/Algorithms/minHash_Jaccard_Similarity.php)

tf-idf weight (/Algorithms/tf_idf_term_frequency-inverse_document_frequency.php)

Natural Language Processing (NLP): Sentiment Analysis I (IMDb & bag-of-words) (/Algorithms/Machine_Learning/NLP/Sentiment_Analysis_I.php)

Natural Language Processing (NLP): Sentiment Analysis II (tokenization, stemming, and stop words) (/Algorithms/Machine_Learning/NLP/Sentiment_Analysis_II.php)

Natural Language Processing (NLP): Sentiment Analysis III (training & cross validation) (/Algorithms/Machine_Learning/NLP/Sentiment_Analysis_III.php)

Natural Language Processing (NLP): Sentiment Analysis IV (out-of-core) (/Algorithms/Machine_Learning/NLP/Sentiment_Analysis_IV.php)

Locality-Sensitive Hashing (LSH) using Cosine Distance (Cosine Similarity) (/Algorithms/Locality_Sensitive_Hashing/LSH_using_Cosine_Distance.php)

ARTIFICIAL NEURAL NETWORKS (ANN)

1. Introduction (/python/scikit-learn/Artificial-Neural-Network-ANN-1-Introduction.php)

2. Forward Propagation (/python/scikit-learn/Artificial-Neural-Network-ANN-2-Forward-Propagation.php)

Note that choosing the number of principal components has to be determined from a trade-off between computational efficiency and the performance of the classifier, however, we only chose two eigenvectors for the demonstration purpose.

```
w= np.hstack((eigen_pairs[0][1][:, np.newaxis], eigen_pairs[1][1][:, np.newaxis]))
w.shape

(13, 2)

w

array([[ 0.14669811,  0.50417079],
       [-0.24224554,  0.24216889],
       [-0.02993442,  0.28698484],
       [-0.25519002, -0.06468718],
       [ 0.12079772,  0.22995385],
       [ 0.38934455,  0.09363991],
       [ 0.42326486,  0.01088622],
       [-0.30634956,  0.01870216],
       [ 0.30572219,  0.03040352],
       [-0.09869191,  0.54527081],
       [ 0.30032535, -0.27924322],
       [ 0.36821154, -0.174365  ],
       [ 0.29259713,  0.36315461]])
```

Now we've created a 13×2 **projection matrix \mathbf{W}** from the top two eigenvectors.

Using the projection matrix, we can now transform a sample \mathbf{x} (represented as 1×13 row vector) onto the PCA subspace obtaining \mathbf{x}' which is a 2-D sample vector consisting of two new features:

$$\mathbf{x}' = \mathbf{x}\mathbf{W}$$

```
X_train_std[0]

array([[ 0.91083058, -0.46259897, -0.01142613, -0.82067872,  0.06241693,
         0.58820446,  0.93565436, -0.7619138 ,  0.13007174, -0.51238741,
         0.65706596,  1.94354495,  0.93700997]])

X_train_std[0].dot(w)

array([ 2.59891628,  0.00484089])
```

In the same way, we can transform the entire 124×13 training dataset onto the two principal components by calculating the matrix dot product:

$$\mathbf{X}' = \mathbf{X}\mathbf{W}$$

```
X_train_pca = X_train_std.dot(w)
X_train_std.shape, w.shape, X_train_pca.shape

((124, 13), (13, 2), (124, 2))
```

Finally, it's time to visualize the transformed Wine training set, now stored as an 124×2 matrix, in a two-dimensional scatterplot:

3. Gradient Descent (/python/scikit-learn/Artificial-Neural-Network-ANN-3-Gradient-Descent.php)

4. Backpropagation of Errors (/python/scikit-learn/Artificial-Neural-Network-ANN-4-Backpropagation.php)

5. Checking gradient (/python/scikit-learn/Artificial-Neural-Network-ANN-5-Checking-Gradient.php)

6. Training via BFGS (/python/scikit-learn/Artificial-Neural-Network-ANN-6-Training-via-BFGS-Broyden-Fletcher-Goldfarb-Shanno-algorithm-a-variant-of-gradient-descent.php)


7. Overfitting & Regularization (/python/scikit-learn/Artificial-Neural-Network-ANN-7-Overfitting-Regularization.php)

8 - Deep Learning I : Image Recognition (Image uploading) (/python/scikit-learn/Artificial-Neural-Network-ANN-8-Deep-Learning-1-Image-Recognition-Image-Uploading.php)

9 - Deep Learning II : Image Recognition (Image classification) (/python/scikit-learn/Artificial-Neural-Network-ANN-9-Deep-Learning-2-Image-Recognition-Image-Classification.php)

10 - Deep Learning III : Deep Learning III : Theano, TensorFlow, and Keras (/python/scikit-learn/Artificial-Neural-Network-ANN-10-Deep-Learning-3-Theano-TensorFlow-Keras.php)

```
colors = ['r', 'b', 'g']
markers = ['s', 'x', 'o']
for l, c, m in zip(np.unique(y_train), colors, markers):
    scatter(X_train_pca[y_train==l, 0], X_train_pca[y_train==l, 1],
            c=c, label=l, marker=m)
xlabel('PC 1')
ylabel('PC 2')
legend(loc='lower left')
show()
```



We we can see from the plot, the data is more spread along the xx -axis which is the first principal component than the yy -axis which is the second principal component.

Though this is consistent with the explained variance ratio plot that we created in the previous subsection, we can intuitively see that a linear classifier will likely be able to separate the classes well.

One more thing to remind once again:
Although we encoded the class labels information for the purpose of illustration in the preceding scatter plot, we have to keep in mind that **PCA is an unsupervised** technique that doesn't use class label information.

PCA in scikit-learn

In this section we want to learn how to use the PCA class implemented in scikit-learn.

PCA is another one of a scikit-learn's transformer classes, where we first fit the model using the training data before we transform both the training data and the test data using the same model parameters.

Let's use the PCA from scikit-learn on the Wine training dataset, and classify the transformed samples via logistic regression.

To visualize the decision regions, we'll use the following code:

Python tutorial

Python Home
(/python/pytut.php)

Introduction
(/python/python_introduction.php)

Running Python Programs (os, sys, import)
(/python/python_running.php)

Modules and IDLE (Import, Reload, exec)
(/python/python_modules_idle.php)

Object Types - Numbers, Strings, and None
(/python/python_numbers_strings.php)

Strings - Escape Sequence, Raw


```
from matplotlib.colors import ListedColormap


def plot_decision_regions(X, y, classifier, resolution=0.02):

    # setup marker generator and color map
    markers = ('s', 'x', 'o', '^', 'v')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    # plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    Z = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    Z = Z.reshape(xx1.shape)
    contourf(xx1, xx2, Z, alpha=0.4, cmap=cmap)
    xlim(xx1.min(), xx1.max())
    ylim(xx2.min(), xx2.max())

    # plot class samples
    for idx, cl in enumerate(np.unique(y)):
        scatter(x=X[y == cl, 0], y=X[y == cl, 1], alpha=0.8, c=cmap(idx),
               marker=markers[idx], label=cl)
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
lr = LogisticRegression()
X_train_pca = pca.fit_transform(X_train_std)
X_test_pca = pca.transform(X_test_std)
lr.fit(X_train_pca, y_train)
plot_decision_regions(X_train_pca, y_train, classifier=lr)
xlabel('PC1')
ylabel('PC2')
legend(loc='lower left')
show()
```



From the picture above, we can see the decision regions for the training model reduced to the two principal component axes.

Note: Covariance-matrix

This section explains the core concept of covariance matrix (ref Statistics 101: The Covariance Matrix (<https://www.youtube.com/watch?v=locZabK4Als>)).

Here is the data we're going to use:

String, and Slicing
(/python/python_strings.php)

Strings - Methods
(/python/python_strings_methods.php)

Formatting Strings -
expressions and method calls
(/python/python_string_formatting.php)

Files and os.path
(/python/python_files.php)

Traversing directories
recursively
(/python/python_traversing_directories.php)

Subprocess Module
(/python/python_subprocess_module.php)

Regular Expressions with
Python
(/python/python_regularExpressions.php)

Object Types - Lists
(/python/python_lists.php)

Object Types - Dictionaries and
Tuples
(/python/python_dictionaries_tuples.php)

Functions def, *args, **kwargs
(/python/python_functions_def_args_kwargs.php)

Functions lambda
(/python/python_functions_lambda.php)

Built-in Functions
(/python/python_functions_built_in.php)

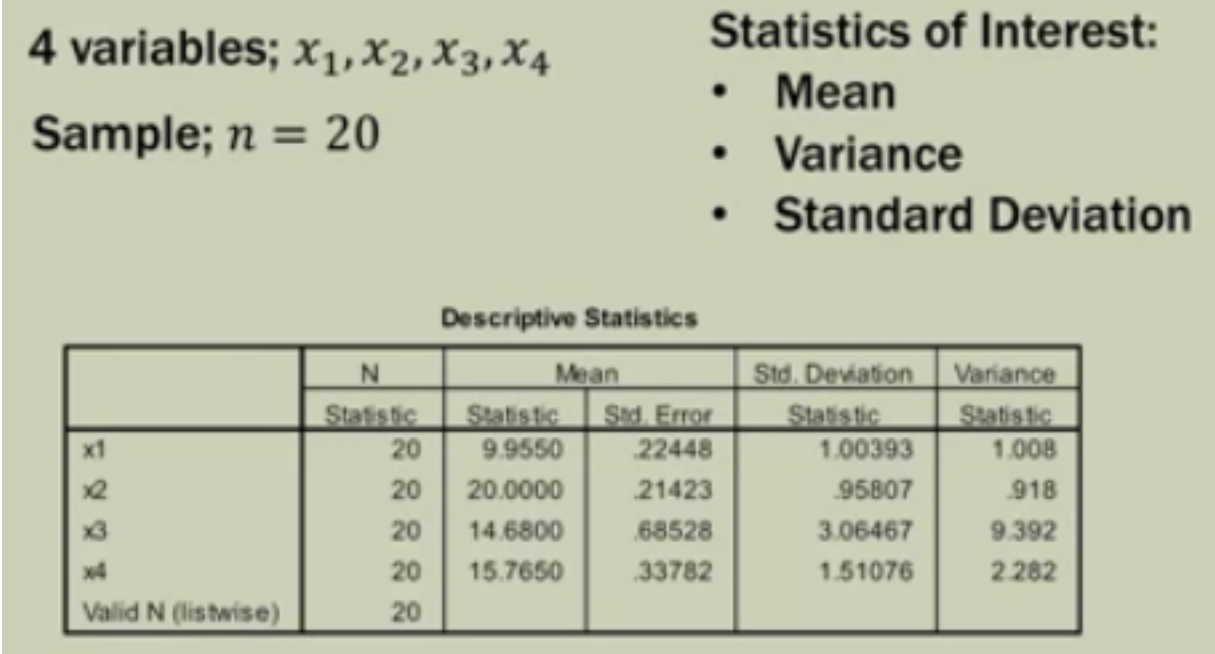
map, filter, and reduce
(/python/python_fncls_map_filter_reduce.php)

Decorators
(/python/python_decorators.php)

List Comprehension
(/python/python_list_comprehension.php)

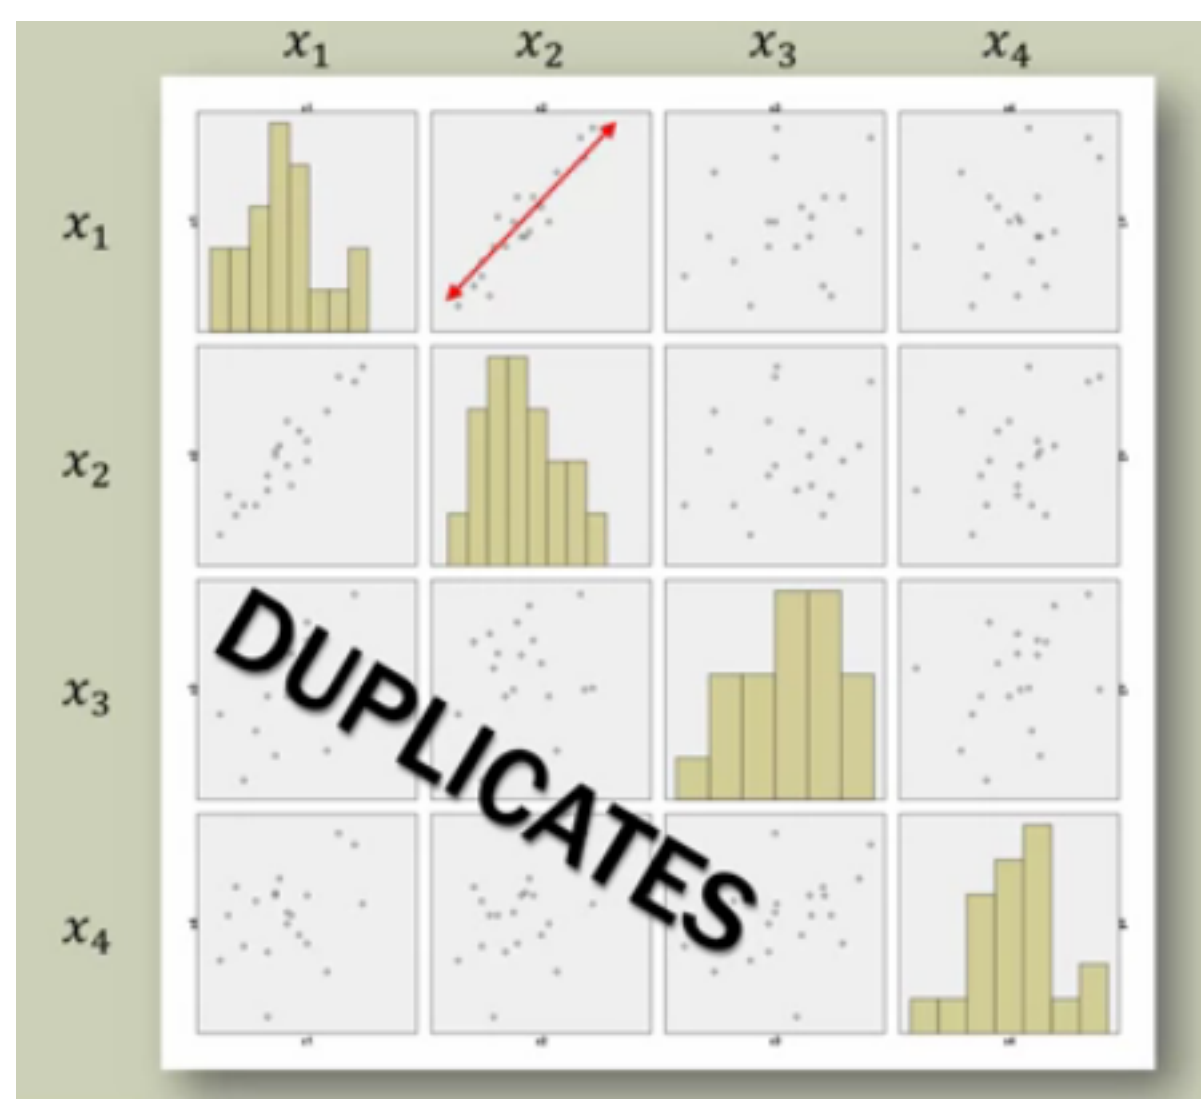
Sets (union/intersection) and
itertools - Jaccard coefficient
and shingling to check
plagiarism
(/python/python_sets_union_intersection.php)

Hashing (Hash tables and
hashlib)
(/python/python_hash_tables_hashlib.php)




- Dictionary Comprehension with zip (/python/python_dictionary_comprehension.php)
- The yield keyword (/python/python_function_with_yield.php)
- Generator Functions and Expressions (/python/python_generators.php)
- generator.send() method (/python/python_function_with_yield.php)
- Iterators (/python/python_iterators.php)
- Classes and Instances (__init__, __call__, etc.) (/python/python_classes_instances.php)
- if __name__ == '__main__' (/python/python_if_name_equals_main.php)
- argparse (/python/python_argparse.php)
- Exceptions (/python/python_try_except_finally.php)
- @static method vs class method (/python/python_differences_between_static_method_and_class_method.php)
- Private attributes and private methods (/python/python_private_attributes_and_methods.php)
- bits, bytes, bitstring, and constBitStream (/python/python_bits_bytes_bitstring_and_constbitstream.php)
- json.dump(s) and json.load(s) (/python/python-json-dumps-loads-file-read-write.php)
- Python Object Serialization - pickle and json (/python/python_serialization_pickle_and_json.php)
- Python Object Serialization - yaml and json (/python/python_yaml_json_serialization.php)
- Priority queue and heap queue data structure (/python/python_PriorityQueue_and_HeapQueue_data_structures.php)

Here is the scatter plots of each variable against the others:



So, we get the following covariance matrix



As we can see from the picture above, the diagonal of a covariance matrix gives us the two kinds of variances: the variance of each variables against others and the covariance with itself. So, the off-diagonal entries are the covariance between pair of variables.

	x_1	x_2	x_3	x_4
x_1	$Var(x_1)$	$Cov(x_1, x_2)$	$Cov(x_1, x_3)$	$Cov(x_1, x_4)$
x_2		$Var(x_2)$	$Cov(x_2, x_3)$	$Cov(x_2, x_4)$
x_3			$Var(x_3)$	$Cov(x_3, x_4)$
x_4				$Var(x_4)$

Github repository

Source is available from bogotobogo-Machine-Learning (<https://github.com/Einsteinish/bogotobogo-Machine-Learning.git>).

Next:

Data Compression via Dimensionality Reduction II - Linear Discriminant Analysis (LDA) ([/python/scikit-learn/scikit_machine_learning_Data_Compresssion_via_Dimensionality_Reduction_2_Linear_Discriminant_Analysis.php](#))

Machine Learning with scikit-learn

scikit-learn installation ([/python/scikit-learn/scikit-learn_install.php](#))

scikit-learn : Features and feature extraction - iris dataset ([/python/scikit-learn/scikit_machine_learning_features_extraction.php](#))

scikit-learn : Machine Learning Quick Preview ([/python/scikit-learn/scikit_machine_learning_quick_preview.php](#))

scikit-learn : Data Preprocessing I - Missing / Categorical data ([/python/scikit-learn/scikit_machine_learning_Data_Preprocessing-Missing-Data-Categorical-Data.php](#))

scikit-learn : Data Preprocessing II - Partitioning a dataset / Feature scaling / Feature Selection / Regularization ([/python/scikit-learn/scikit_machine_learning_Data_Preprocessing-II-Datasets-Partitioning-Feature-scaling-Feature-Selection-Regularization.php](#))

scikit-learn : Data Preprocessing III - Dimensionality reduction vis Sequential feature selection / Assessing feature importance via random forests ([/python/scikit-learn/scikit_machine_learning_Data_Preprocessing-III-Dimensionality-reduction-via-Sequential-feature-selection-Assessing-feature-importance-via-random-forests.php](#))

Data Compression via Dimensionality Reduction I - Principal component analysis (PCA) ([/python/scikit-learn/scikit_machine_learning_Data_Compresssion_via_Dimensionality_Reduction_1_Principal_component_analysis_PCA.php](#))

scikit-learn : Data Compression via Dimensionality Reduction II - Linear Discriminant Analysis (LDA) ([/python/scikit-learn/scikit_machine_learning_Data_Compresssion_via_Dimensionality_Reduction_2_Linear_Discriminant_Analysis.php](#))

scikit-learn : Data Compression via Dimensionality Reduction III - Nonlinear mappings via kernel

Graph data structure ([/python/python_graph_data_struct.php](#))

Dijkstra's shortest path algorithm ([/python/python_Dijkstras_Shortest_Path.php](#))

Prim's spanning tree algorithm ([/python/python_Prims_Spanning_Tree.php](#))

Closure ([/python/python_closure.php](#))
Functional programming in Python ([/python/python_functional_programming.php](#))

Remote running a local file using ssh ([/python/python_ssh_remote_run.php](#))

SQLite 3 - A. Connecting to DB, create/drop table, and insert data into a table ([/python/python_sqlite_connect.php](#))

SQLite 3 - B. Selecting, updating and deleting data ([/python/python_sqlite_select_update_delete.php](#))

MongoDB with PyMongo I - Installing MongoDB ... ([/python/MongoDB_PyMongo/python_mongodb_installation.php](#))

Python HTTP Web Services - urllib, httplib2 ([/python/python_http_web_services.php](#))

Web scraping with Selenium for checking domain availability ([/python/python_Web_scraping_with_Selenium.php](#))

REST API : Http Requests for Humans with Flask ([/python/python-REST-API-Http-Requests-for-Humans-with-Flask.php](#))

Blog app with Tornado ([/python/Tornado/Python_Tornado_blog_app.php](#))

Multithreading ... ([/python/Multithread/python_multithreading.php](#))
Python Network Programming I - Basic Server / Client : A

principal component (KPCA) analysis (/python/scikit-learn/scikit_machine_learning_Data_Compresssion_via_Dimensionality_Reduction_3-nonlinear-mappings-via-kernel-principal-component-analysis.php)	Basics (/python/python_network_programming.php)
scikit-learn : Logistic Regression, Overfitting & regularization (/python/scikit-learn/scikit_learn_logistic_regression.php)	Python Network Programming I - Basic Server / Client : B File Transfer (/python/python_network_programming.php)
scikit-learn : Supervised Learning & Unsupervised Learning - e.g. Unsupervised PCA dimensionality reduction with iris dataset (/python/scikit-learn/scikit_machine_learning_Supervised_Learning_Unsupervised_Learning.php)	Python Network Programming II - Chat Server / Client (/python/python_network_programming.php)
scikit-learn : Unsupervised_Learning - KMeans clustering with iris dataset (/python/scikit-learn/scikit_machine_learning_Unsupervised_Learning_Clustering.php)	Python Network Programming III - Echo Server using socketserver network framework (/python/python_network_programming.php)
scikit-learn : Linearly Separable Data - Linear Model & (Gaussian) radial basis function kernel (RBF kernel) (/python/scikit-learn/scikit_machine_learning_Linearly_Separable_NonLinearly_RBF_Separable_Data_SVM_GUI.php)	Python Network Programming IV - Async From a Request Handling : ThreadingMixIn and ForkingMixIn (/python/python_network_programming.php)
scikit-learn : Decision Tree Learning I - Entropy, Gini, and Information Gain (/python/scikit-learn/scikt_machine_learning_Decision_Tree_Learning_Informatioin_Gain_IG_Impurity_Entropy_Gini_Classification.php)	Python Interview Questions I (/python/python_interview_questions.php)
scikit-learn : Decision Tree Learning II - Constructing the Decision Tree (/python/scikit-learn/scikit_machine_learning_Constructing_Decision_Tree_Learning_Information_Gain_IG_Impurity_Entropy_Gini_Classification.php)	Python Interview Questions II (/python/python_interview_questions.php)
scikit-learn : Random Decision Forests Classification (/python/scikit-learn/scikit_machine_learning_Random_Decision_Forests_Ensemble_Learning_Classification.php)	Python Interview Questions III (/python/python_interview_questions.php)
scikit-learn : Support Vector Machines (SVM) (/python/scikit-learn/scikit_machine_learning_Support_Vector_Machines_SVM.php)	Python Interview Questions IV (/python/python_interview_questions.php)
scikit-learn : Support Vector Machines (SVM) II (/python/scikit-learn/scikit_machine_learning_Support_Vector_Machines_SVM_2.php)	Python Interview Questions V (/python/python_interview_questions.php)
Flask with Embedded Machine Learning I : Serializing with pickle and DB setup (/python/Flask/Python_Flask_Embedding_Machine_Learning_1.php)	Image processing with Python image library Pillow (/python/python_image_processing.php)
Flask with Embedded Machine Learning II : Basic Flask App (/python/Flask/Python_Flask_Embedding_Machine_Learning_2.php)	Python and C++ with SIP (/python/python_cpp_sip.php)
Flask with Embedded Machine Learning III : Embedding Classifier (/python/Flask/Python_Flask_Embedding_Machine_Learning_3.php)	PyDev with Eclipse (/python/pydev_eclipse_plugin_installation.php)
Flask with Embedded Machine Learning IV : Deploy (/python/Flask/Python_Flask_Embedding_Machine_Learning_4.php)	Matplotlib (/python/python_matplotlib.php)
Flask with Embedded Machine Learning V : Updating the classifier (/python/Flask/Python_Flask_Embedding_Machine_Learning_5.php)	Redis with Python (/python/python_redis_with_python.php)
scikit-learn : Sample of a spam comment filter using SVM - classifying a good one or a bad one (/python/scikit-learn/scikit_learn_Support_Vector_Machines_SVM_spam_filtermachine_learning_.php)	NumPy array basics A (/python/python_numpy_array_tutorial.php)
<hr/>	
<h2>MACHINE LEARNING ALGORITHMS AND CONCEPTS</h2>	
Batch gradient descent algorithm (/python/python_numpy_batch_gradient_descent_algorithm.php)	

Single Layer Neural Network - Perceptron model on the Iris dataset using Heaviside step activation function (/python/scikit-learn/Perceptron_Model_with_Iris_DataSet.php)

Batch gradient descent versus stochastic gradient descent (/python/scikit-learn/scikit-learn_batch-gradient-descent-versus-stochastic-gradient-descent.php)

Single Layer Neural Network - Adaptive Linear Neuron using linear (identity) activation function with batch gradient descent method (/python/scikit-learn/Single-Layer-Neural-Network-Adaptive-Linear-Neuron.php)

Single Layer Neural Network : Adaptive Linear Neuron using linear (identity) activation function with stochastic gradient descent (SGD) (/python/scikit-learn/Single-Layer-Neural-Network-Adaptive-Linear-Neuron-with-Stochastic-Gradient-Descent.php)

Logistic Regression (/python/scikit-learn/logistic_regression.php)

VC (Vapnik-Chervonenkis) Dimension and Shatter (/python/scikit-learn/scikit_machine_learning_VC_Dimension_Shatter.php)

Bias-variance tradeoff (/python/scikit-learn/scikit_machine_learning_Bias-variance-Tradeoff.php)

Maximum Likelihood Estimation (MLE) (/python/scikit-learn/Maximum-Likelyhood-Estimation-MLE.php)

Neural Networks with backpropagation for XOR using one hidden layer (/python/python_Neural_Networks_Backpropagation_for_XOR_using_one_hidden_layer.php)

minHash (/Algorithms/minHash_Jaccard_Similarity_Locality_sensitive_hashing_LSH.php)

tf-idf weight (/Algorithms/tf_idf_term_frequency_inverse_document_frequency_NLP_Natural_Language_Processing.php)

Natural Language Processing (NLP): Sentiment Analysis I (IMDb & bag-of-words) (/Algorithms/Machine_Learning_NLP_Sentiment_Analysis_1.php)

Natural Language Processing (NLP): Sentiment Analysis II (tokenization, stemming, and stop words) (/Algorithms/Machine_Learning_NLP_Sentiment_Analysis_2.php)

Natural Language Processing (NLP): Sentiment Analysis III (training & cross validation) (/Algorithms/Machine_Learning_NLP_Sentiment_Analysis_3.php)

Natural Language Processing (NLP): Sentiment Analysis IV (out-of-core) (/Algorithms/Machine_Learning_NLP_Sentiment_Analysis_4.php)

Locality-Sensitive Hashing (LSH) using Cosine Distance (Cosine Similarity) (/Algorithms/Locality_Sensitive_Hashing_LSH_using_Cosine_Distance_Similarity.php)

ARTIFICIAL NEURAL NETWORKS (ANN)

[Note] Sources are available at Github - Jupyter notebook files (https://github.com/Einsteinish/Artificial-Neural-Networks-with-Jupyter.git)

1. Introduction (/python/scikit-learn/Artificial-Neural-Network-ANN-1-Introduction.php)

NumPy Matrix and Linear Algebra (/python/python_numpy_matrix_and_linear_algebra.php)

Pandas with NumPy and Matplotlib (/python/python_Pandas_NumPy_and_Matplotlib.php)

Celluar Automata (/python/python_cellular_automata.php)

Batch gradient descent algorithm (/python/python_numpy_batch_gradient_descent_algorithm.php)

Longest Common Substring Algorithm (/python/python_longest_common_substring_algorithm.php)

Python Unit Test - TDD using unittest.TestCase class (/python/python_unit_testing_python_unittest_TestCase_class.php)

Simple tool - Google page ranking by keywords (/python/python_site_page_ranking_by_keywords.php)

Google App Hello World (/python/GoogleApp/python_GoogleAppHelloWorld.php)

Google App webapp2 and WSGI (/python/GoogleApp/python_GoogleAppWebapp2andWSGI.php)

Uploading Google App Hello World (/python/GoogleApp/python_GoogleAppHelloWorld.php)

Python 2 vs Python 3 (/python/python_differences_Python2vsPython3.php)

virtualenv and virtualenvwrapper (/python/python_virtualenv_virtualenvwrapper.php)

Uploading a big file to AWS S3 using boto module (/DevOps/AWS/aws_S3_uploading_a_big_file_to_AWS_S3_using_boto_module.php)

Scheduled stopping and starting an AWS instance (/DevOps/AWS/aws_stopping_and_starting_an_AWS_instance.php)

Cloudera CDH5 - Scheduled stopping and starting services (/Hadoop/BigData_hadoop_CDH5_Scheduled_stopping_and_starting_services.php)

Removing Cloud Files -

2. Forward Propagation (/python/scikit-learn/Artificial-Neural-Network-ANN-2-Forward-Propagation.php)	Rackspace API with curl and subprocess (/python/python_Rackspace_API/)
3. Gradient Descent (/python/scikit-learn/Artificial-Neural-Network-ANN-3-Gradient-Descent.php)	Checking if a process is running/hanging and stop/run a scheduled task on Windows (/python/python-Windows-Check-if-a-Process-is-Running-Hanging-Schtasks-Run-Stop.php)
4. Backpropagation of Errors (/python/scikit-learn/Artificial-Neural-Network-ANN-4-Backpropagation.php)	Apache Spark 1.3 with PySpark (Spark Python API) Shell (/Hadoop/BigData_hadoop_ApacheSpark/)
5. Checking gradient (/python/scikit-learn/Artificial-Neural-Network-ANN-5-Checking-Gradient.php)	Apache Spark 1.2 Streaming (/Hadoop/BigData_hadoop_ApacheSpark/)
6. Training via BFGS (/python/scikit-learn/Artificial-Neural-Network-ANN-6-Training-via-BFGS-Broyden-Fletcher-Goldfarb-Shanno-algorithm-a-variant-of-gradient-descent.php)	bottle 0.12.7 - Fast and simple WSGI-micro framework for small web-applications ... (/python/Bottle/Python_Bottle_Framework/)
7. Overfitting & Regularization (/python/scikit-learn/Artificial-Neural-Network-ANN-7-Overfitting-Regularization.php)	Flask app with Apache WSGI on Ubuntu14/CentOS7 ... (/python/Flask/Python_Flask_Application/)
8. Deep Learning I : Image Recognition (Image uploading) (/python/scikit-learn/Artificial-Neural-Network-ANN-8-Deep-Learning-1-Image-Recognition-Image-Uploading.php)	Selenium WebDriver (/python/python_Selenium_Webdriver/)
9. Deep Learning II : Image Recognition (Image classification) (/python/scikit-learn/Artificial-Neural-Network-ANN-9-Deep-Learning-2-Image-Recognition-Image-Classification.php)	Fabric - streamlining the use of SSH for application deployment (/python/Fabric/python_Fabric_Application/)
10 - Deep Learning III : Deep Learning III : Theano, TensorFlow, and Keras (/python/scikit-learn/Artificial-Neural-Network-ANN-10-Deep-Learning-3-Theano-TensorFlow-Keras.php)	Ansible Quick Preview - Setting up web servers with Nginx, configure enviroments, and deploy an App (/DevOps/Ansible/Ansible_Setting_up_Web_Servers/)
	Neural Networks with backpropagation for XOR using one hidden layer (/python/python_Neural_Networks/)
	NLP - NLTK (Natural Language Toolkit) ... (/python/NLTK/NLTK_install.php)
	RabbitMQ(Message broker server) and Celery(Task queue) ... (/python/RabbitMQ_Celery/python_RabbitMQ_Celery/)
	OpenCV3 and Matplotlib ... (/python/OpenCV_Python/python_OpenCV3_Matplotlib/)

Simple tool - Concatenating
slides using FFmpeg ...
(/FFMpeg/ffmpeg_fade_in_fade_

iPython - Signal Processing
with NumPy
(/python/OpenCV_Python/pytho

iPython and Jupyter - Install
Jupyter, iPython Notebook,
drawing with Matplotlib, and
publishing it to Github
(/python/IPython/IPython_Jupyte

iPython and Jupyter Notebook
with Embedded D3.js
(/python/IPython/iPython_Jupyte

Downloading YouTube videos
using youtube-dl embedded
with Python
(/VideoStreaming/YouTube/yout
dl-embedding.php)

Machine Learning : scikit-learn
... (/python/scikit-
learn/scikit_machine_learning_S

Django 1.6/1.8 Web
Framework ...
(/python/Django/Python_Django

CONTACT

BogoToBogo
contactus@bogotobogo.com (mailto:#)

FOLLOW BOGOTOBOGO

f (https://www.facebook.com/KHongSanFrancisco) **t**
(https://twitter.com/KHongTwit) **g**⁺
(https://plus.google.com/u/0/+KHongSanFrancisco/posts)

ABOUT US (/ABOUT_US.PHP)

contactus@bogotobogo.com (mailto:#)

Golden Gate Ave, San Francisco, CA 94115

Loading [Mathjax]/jax/output/CommonHTML/jax.js